

Project 3 - DS8002

Rafik Matta

Exploratory Data Analysis

```
library(readr)
library("MASS")
library(corrplot)
```

```
## corrplot 0.84 loaded
```

```
library(plotly)
```

```
## Loading required package: ggplot2
```

```
##
```

```
## Attaching package: 'plotly'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
##     last_plot
```

```
## The following object is masked from 'package:MASS':
```

```
##
```

```
##     select
```

```
## The following object is masked from 'package:stats':
```

```
##
```

```
##     filter
```

```
## The following object is masked from 'package:graphics':
```

```
##
```

```
##     layout
```

```
winequality_red <- read_delim("C:/Users/rafik/Google Drive/Education/ds8002/projects/project3/winequality_red.csv",  
                             ";", escape_double = FALSE, trim_ws = TRUE)
```

```
## Parsed with column specification:
```

```
## cols(
```

```
##   `fixed acidity` = col_double(),
```

```
##   `volatile acidity` = col_double(),
```

```
##   `citric acid` = col_double(),
```

```
##   `residual sugar` = col_double(),
```

```
##   chlorides = col_double(),
```

```
##   `free sulfur dioxide` = col_double(),
```

```
##   `total sulfur dioxide` = col_integer(),
```

```
##   density = col_double(),
```

```
##   pH = col_double(),
```

```
##   sulphates = col_double(),
```

```
##   alcohol = col_double(),
```

```
##   quality = col_integer()
```

```
## )
```

```
## Warning in rbind(names(probs), probs_f): number of columns of result is not
```

```
## a multiple of vector length (arg 1)
```

```

## Warning: 2 parsing failures.
## row # A tibble: 2 x 5 col      row      col      expected actual expected  <int>
winequality_white <- read_delim("C:/Users/rafik/Google Drive/Education/ds8002/projects/project3/winequa
";", escape_double = FALSE, trim_ws = TRUE)

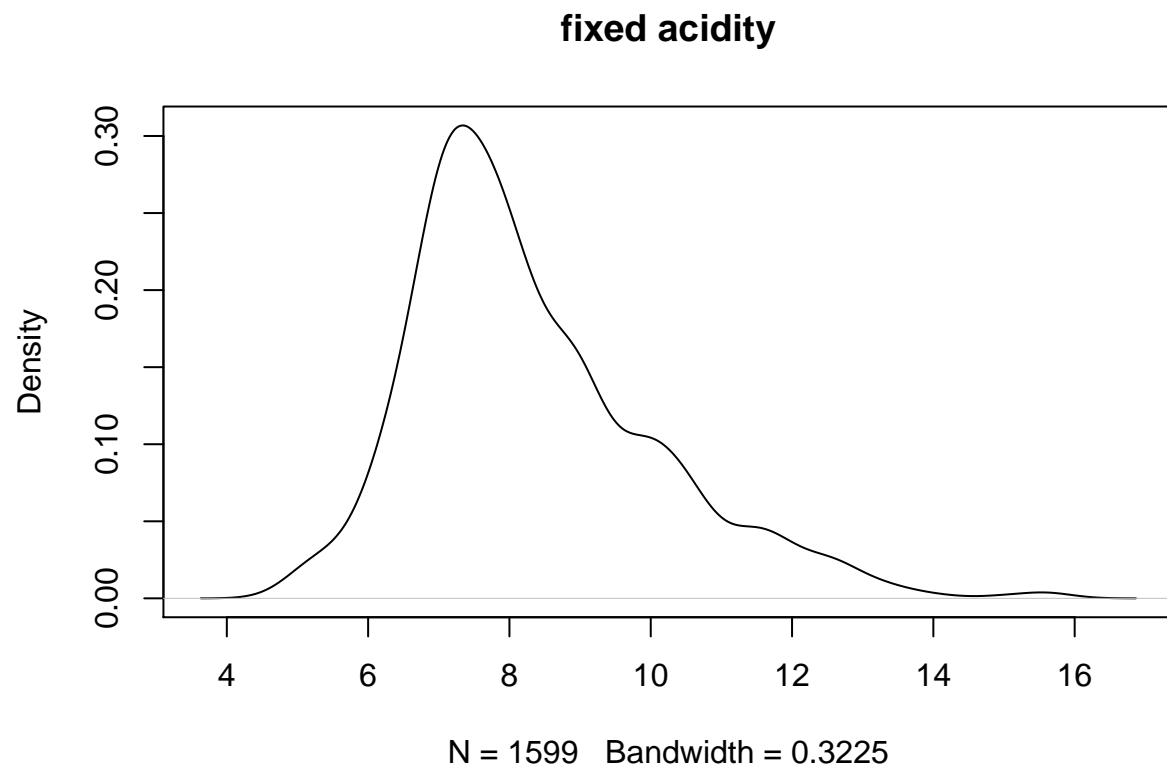
## Parsed with column specification:
## cols(
##   `fixed acidity` = col_double(),
##   `volatile acidity` = col_double(),
##   `citric acid` = col_double(),
##   `residual sugar` = col_double(),
##   chlorides = col_double(),
##   `free sulfur dioxide` = col_double(),
##   `total sulfur dioxide` = col_double(),
##   density = col_double(),
##   pH = col_double(),
##   sulphates = col_double(),
##   alcohol = col_double(),
##   quality = col_integer()
## )


```

EDA Question 1- Variance of each column


Plots for Red DataSet

```
for( i in colnames(winequality_red))  
{  
  d <- density(winequality_red[[i]],na.rm = TRUE)  
  plot(d,main=i)  
}
```






citric acid




N = 1599 Bandwidth = 0.04009

residual sugar




N = 1599 Bandwidth = 0.1075

chlorides




N = 1599 Bandwidth = 0.003072

free sulfur dioxide




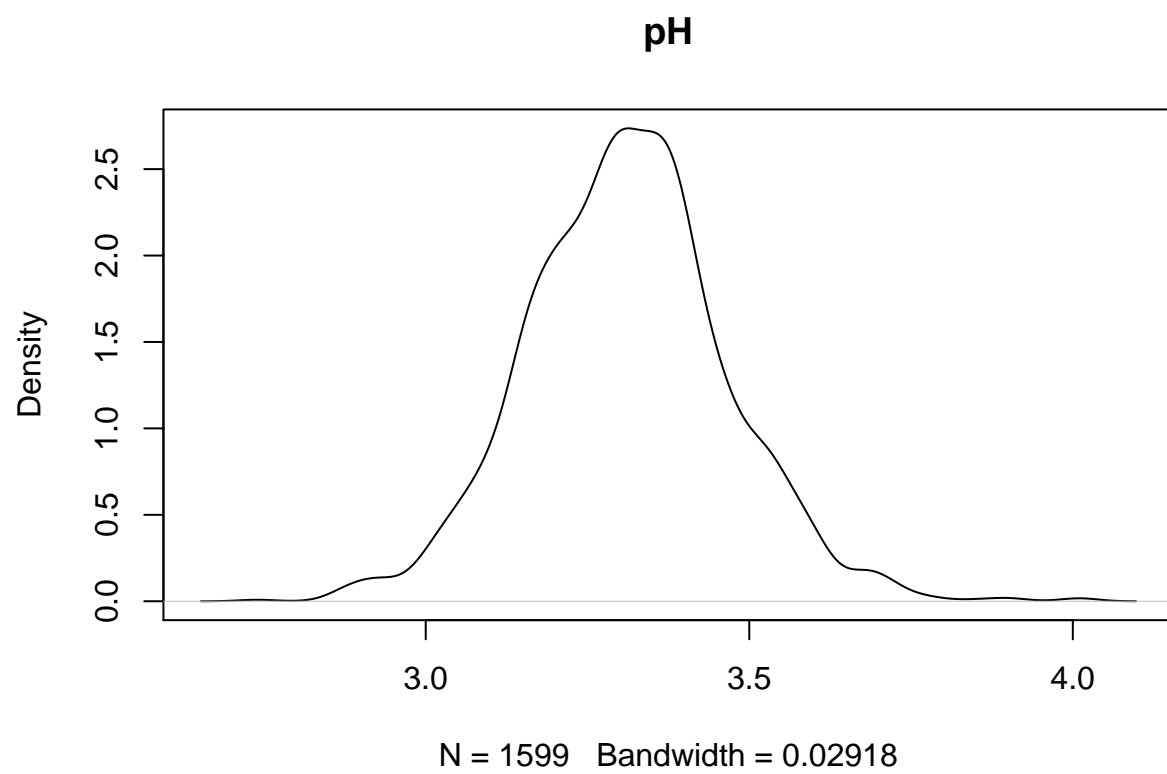
N = 1599 Bandwidth = 2.15

total sulfur dioxide




N = 1597 Bandwidth = 6.145






sulphates

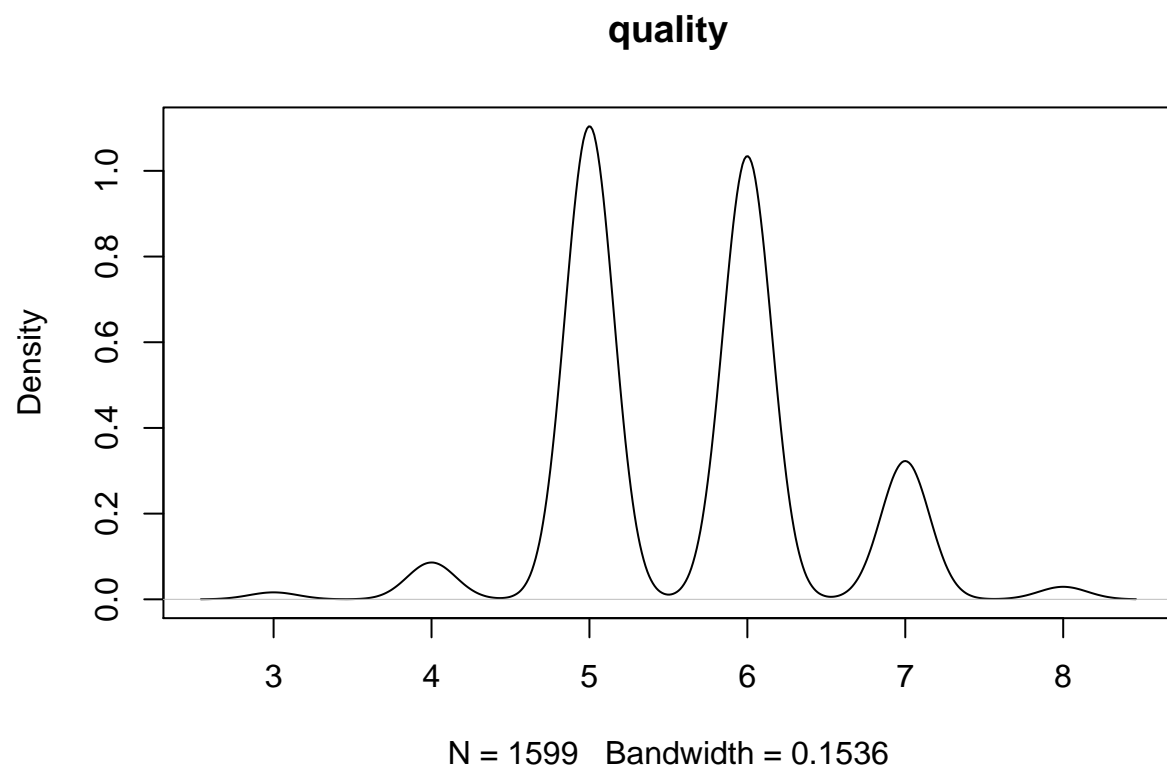


N = 1599 Bandwidth = 0.02765

alcohol




N = 1599 Bandwidth = 0.2193



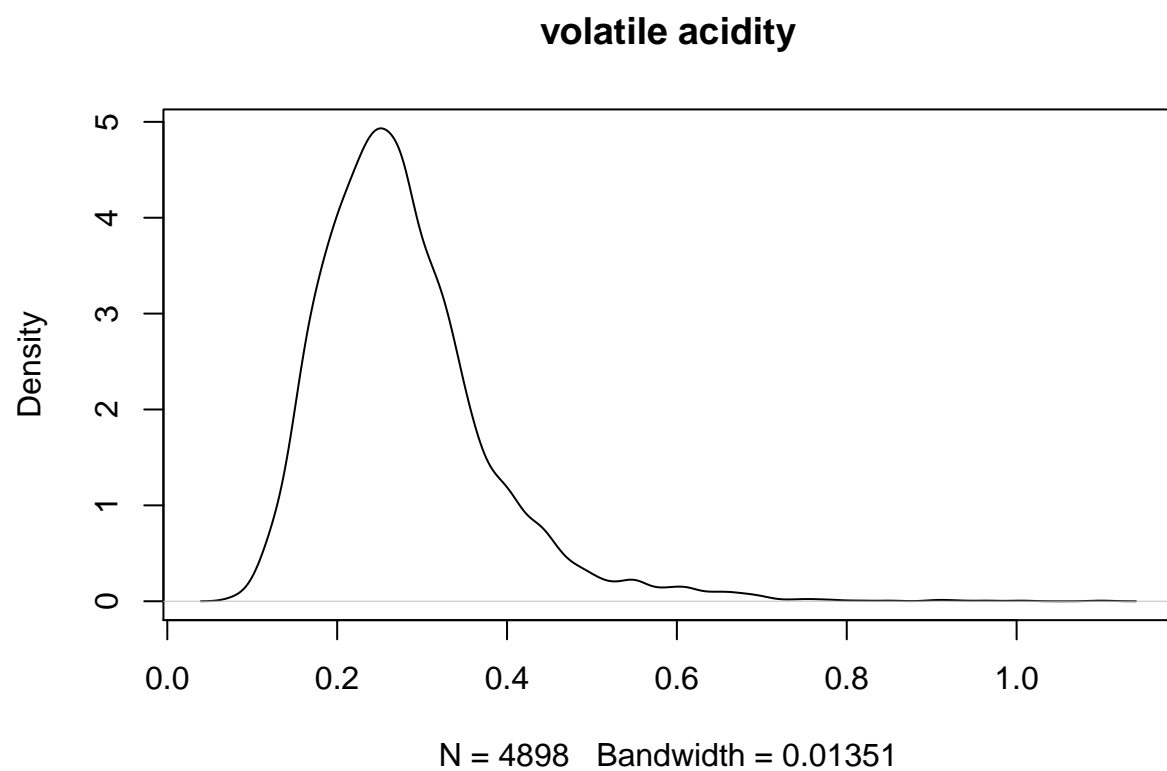
Plots for White DataSet

```
for( i in colnames(winequality_white))  
{  
  d <- density(winequality_white[[i]])  
  plot(d,main=i)  
}
```


fixed acidity



N = 4898 Bandwidth = 0.1228




citric acid




N = 4898 Bandwidth = 0.01473

residual sugar




chlorides

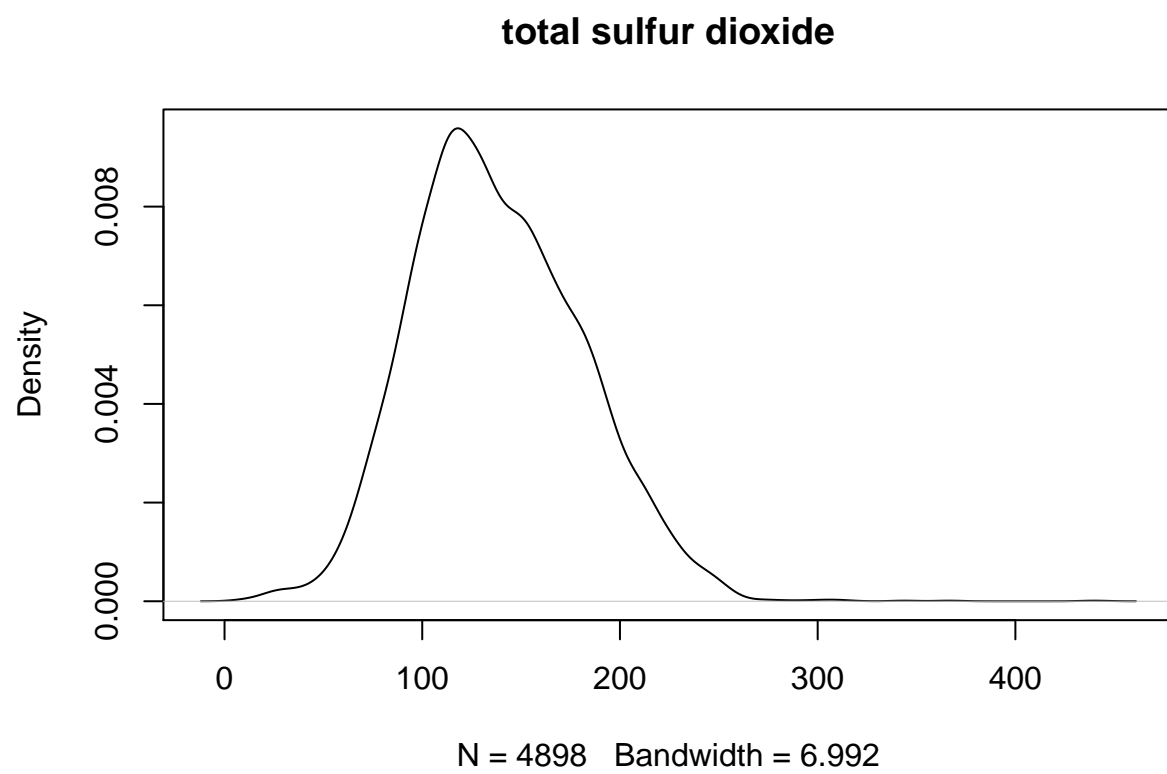



N = 4898 Bandwidth = 0.001719

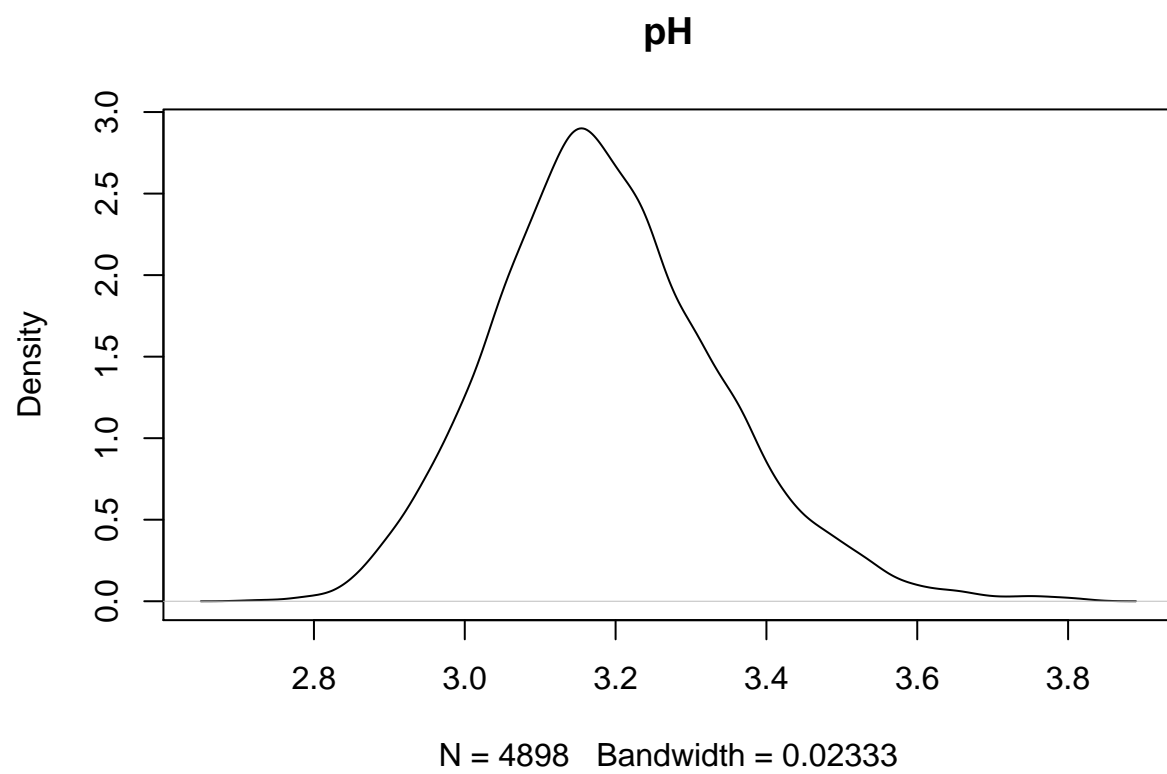
free sulfur dioxide




N = 4898 Bandwidth = 2.798








sulphates

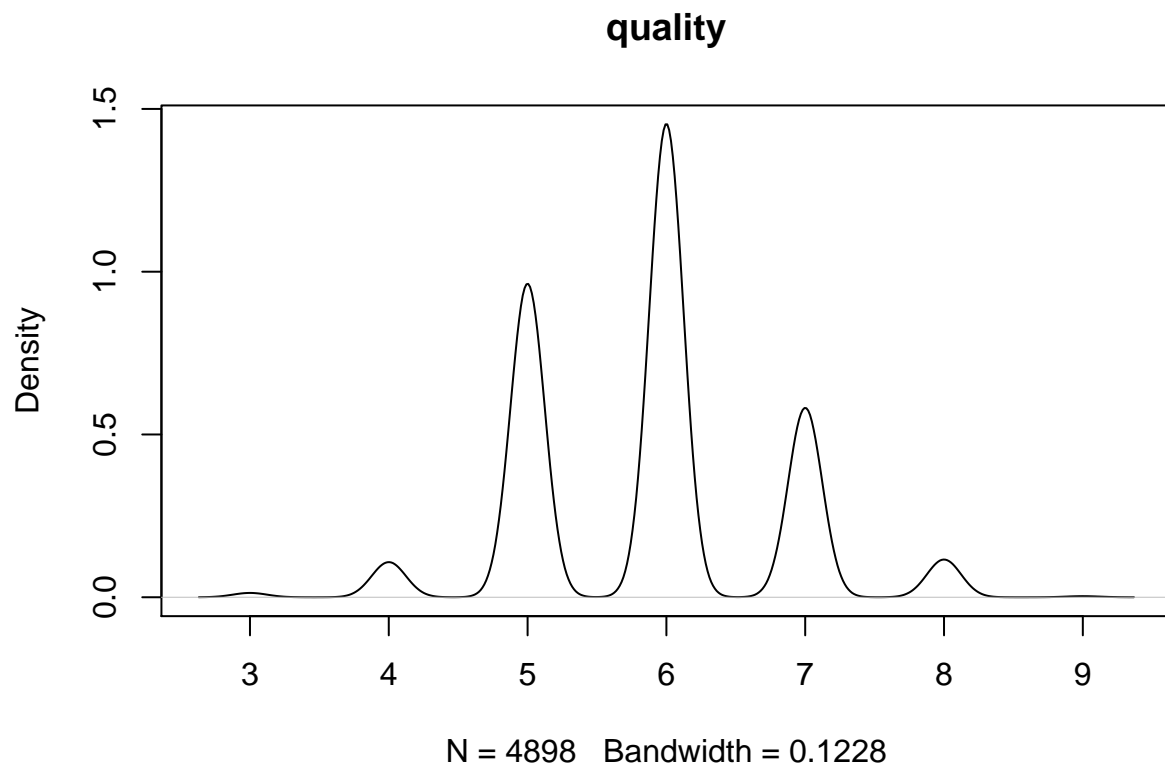


N = 4898 Bandwidth = 0.01719

alcohol



N = 4898 Bandwidth = 0.2025



```
print("RED WINES")

## [1] "RED WINES"
print("=====")

## [1] "===== "

for( i in colnames(winequality_red))
{
  print(i)
  print(var(winequality_red[[i]],na.rm=TRUE))
  print("-----")
}

## [1] "fixed acidity"
## [1] 3.031416
## [1] "-----"
## [1] "volatile acidity"
## [1] 0.03206238
## [1] "-----"
## [1] "citric acid"
## [1] 0.03794748
## [1] "-----"
## [1] "residual sugar"
## [1] 1.987897
## [1] "-----"
## [1] "chlorides"
```

```

## [1] 0.002215143
## [1] "-----"
## [1] "free sulfur dioxide"
## [1] 109.4149
## [1] "-----"
## [1] "total sulfur dioxide"
## [1] 1082.25
## [1] "-----"
## [1] "density"
## [1] 3.562029e-06
## [1] "-----"
## [1] "pH"
## [1] 0.02383518
## [1] "-----"
## [1] "sulphates"
## [1] 0.02873262
## [1] "-----"
## [1] "alcohol"
## [1] 1.135647
## [1] "-----"
## [1] "quality"
## [1] 0.6521684
## [1] "-----"

print("WHITE WINES")

## [1] "WHITE WINES"

print("=====")

## [1] "===== "

for( i in colnames(winequality_red))
{
  print(i)
  print(var(winequality_red[[i]],na.rm=TRUE))
  print("-----")
}


## [1] "fixed acidity"
## [1] 3.031416
## [1] "-----"
## [1] "volatile acidity"
## [1] 0.03206238
## [1] "-----"
## [1] "citric acid"
## [1] 0.03794748
## [1] "-----"
## [1] "residual sugar"
## [1] 1.987897
## [1] "-----"
## [1] "chlorides"
## [1] 0.002215143
## [1] "-----"
## [1] "free sulfur dioxide"
## [1] 109.4149

```


```
## [1] "-----"
## [1] "total sulfur dioxide"
## [1] 1082.25
## [1] "-----"
## [1] "density"
## [1] 3.562029e-06
## [1] "-----"
## [1] "pH"
## [1] 0.02383518
## [1] "-----"
## [1] "sulphates"
## [1] 0.02873262
## [1] "-----"
## [1] "alcohol"
## [1] 1.135647
## [1] "-----"
## [1] "quality"
## [1] 0.6521684
## [1] "-----"
```

EDA Question 2- Variance of each column

```
corrplot(cor(na.omit(winequality_red)),method = 'number')
```



```
corrplot(cor(winequality_white),method = 'number')
```



EDA Question 3- Feature Selection

Having explored both data set's I will pick red wines as it has a more interesting set of correlations between it's features and generally higher variability overall since it's a larger dataset.

Based on the results of the correlation analysis it's clear that the best features with the highest correlation in the data set chosen are the "Volatile Acidity" feature and the "Alcohol" feature. Using these two, I will prepare the data set for analysis.

```
wq_red_df <- data.frame(winequality_red$`volatile acidity`,
                        winequality_red$`alcohol`,
                        winequality_red$`quality`)

data_size = length(wq_red_df$winequality_red..volatile.acidity.)
train_val = .8*data_size
test_val = train_val + 1

response_var_train <- winequality_red$`quality`[1:train_val]
response_var_test <- winequality_red$`quality`[test_val:data_size]

wq_data <- data.frame(rep(1,length(winequality_red$`quality`)),
                      wq_red_df$winequality_red..volatile.acidity.,
                      wq_red_df$winequality_red.alcohol)
wq_data_matrix <- data.matrix(wq_data)

wq_data_matrix_train <- wq_data_matrix[1:train_val,]
wq_data_matrix_test <- wq_data_matrix[test_val:data_size,]
```

Analysis

Analysis Question 1- Multivariate Linear Regression

As per equation (5.35) from the I2ML course text book:

$$r^t = w_0 + w_1x_1^t + w_2x_2^t + \dots + w_dx_d^t + \epsilon$$

Now I train the model to determine the weights of the regression. This is based off of the derived formula from the book (5.39):

```
weights = ginv(t(wq_data_matrix_train) %*% wq_data_matrix_train) %*%  
          t(wq_data_matrix_train) %*% response_var_train
```

Calculating Epsilon:

```
response_on_training = vector()  
  
for(i in 1:length(response_var_train))  
{  
  response_on_training[i] = weights[1] +  
    wq_data_matrix_train[i,2]*weights[2] +  
    wq_data_matrix_train[i,3]*weights[3]  
}  
  
#this is the expected error  
epsilon = 0  
for(i in 1:length(response_var_train))  
{  
  epsilon = epsilon + response_on_training[i] - weights[1] -  
    wq_data_matrix_train[i,2]*weights[2] -  
    wq_data_matrix_train[i,3]*weights[3]  
}  
epsilon = (0.5)*(epsilon^2)
```

An explanation of the vector that is produced. The first value is

$$w_0$$

which is the intercept. The rest of the values are the ordered weights of the variables.

With my weights produced, I can now test my model against the test data:

```
response_test = vector()  
for(i in 1:length(response_var_test))  
{  
  response_test[i] = weights[1] +  
    wq_data_matrix_test[i,2]*weights[2] +  
    wq_data_matrix_test[i,3]*weights[3]  
}
```

Finally, to test the performance of the model, I will calculate the RMSE and the R-Squared:

```
#RMSE  
rmse_of_model = sqrt(sum((response_test-response_var_test)^2))  
  
#R-Squared
```



```

ssr = sum((response_test-response_var_test)^2)
ss_total = sum((response_test-mean(response_test))^2)
rsquared = (ss_total - ssr)/ss_total

print("RMSE:")

## [1] "RMSE:"
print(rmse_of_model)

## [1] 12.08537
print("R-Squared:")

## [1] "R-Squared:"
print(rsquared)

## [1] -1.638753

```

As can be seen by the above values, it's quite clear that this model performs very poorly. As this model does have an intercept as mentioned above, the RSquared value being negative is an indication of an extremely poor fit to the data.

Analysis Question 2- Multivariate Quadratic Regression (Higher Order Polynomial)

I will attempt to see if a quadratic regression with a higher order of polynomial can produce a better result.

As before I will first train the data.

```
wq_data_matrix_train = cbind(wq_data_matrix_train,wq_data_matrix_train[,2]^2)
wq_data_matrix_train = cbind(wq_data_matrix_train,wq_data_matrix_train[,2]*wq_data_matrix_train[,3])
wq_data_matrix_train = cbind(wq_data_matrix_train,wq_data_matrix_train[,3]^2)

weights = ginv(t(wq_data_matrix_train) %*% wq_data_matrix_train) %*%
          t(wq_data_matrix_train) %*% response_var_train
```

Calculating epsilon:

```
response_on_training = vector()

for(i in 1:length(response_var_train))
{
  response_on_training[i] = weights[1] +
    wq_data_matrix_train[i,2]*weights[2] +
    wq_data_matrix_train[i,3]*weights[3] +
    wq_data_matrix_train[i,4]*weights[4] +
    wq_data_matrix_train[i,5]*weights[5] +
    wq_data_matrix_train[i,6]*weights[6]
}

#this is the expected error
epsilon = 0
for(i in 1:length(response_var_train))
{
  epsilon = epsilon + response_on_training[i] - weights[1] -
    wq_data_matrix_train[i,2]*weights[2] -
    wq_data_matrix_train[i,3]*weights[3] -
    wq_data_matrix_train[i,4]*weights[4] -
    wq_data_matrix_train[i,5]*weights[5] -
    wq_data_matrix_train[i,6]*weights[6]
}
epsilon = (0.5)*(epsilon^2)
```

With my weights produced, I can now test my model against the test data:

```
response_test = vector()

wq_data_matrix_test = cbind(wq_data_matrix_test,wq_data_matrix_test[,2]^2)
wq_data_matrix_test = cbind(wq_data_matrix_test,wq_data_matrix_test[,2]*wq_data_matrix_test[,3])
wq_data_matrix_test = cbind(wq_data_matrix_test,wq_data_matrix_test[,3]^2)

for(i in 1:length(response_var_test))
{
  response_test[i] = weights[1] +
    wq_data_matrix_test[i,2]*weights[2] +
    wq_data_matrix_test[i,3]*weights[3] +
    wq_data_matrix_test[i,4]*weights[4] +
    wq_data_matrix_test[i,5]*weights[5] +
    wq_data_matrix_test[i,6]*weights[6]
}
```

```
wq_data_matrix_test[i,6]*weights[6]
}
```

Finally, to test the performance of the model, I will calculate the RMSE and the R-Squared:

```
#RMSE
rmse_of_model = sqrt(sum((response_test-response_var_test)^2))
```

```
#R-Squared
ssr = sum((response_test-response_var_test)^2)
ss_total = sum((response_test-mean(response_test))^2)
rsquared = (ss_total - ssr)/ss_total
```

```
print("RMSE:")
```

```
## [1] "RMSE:"
```

```
print(rmse_of_model)
```

```
## [1] 12.11713
```

```
print("R-Squared:")
```

```
## [1] "R-Squared:"
```

```
print(rsquared)
```

```
## [1] -1.619055
```

Analysis Question 3 - Which order gives best performance

I did not have time to plot or look at any order greater than 2. That being said, an order of two didn't majorily improve performance so I'm not expecting test error to go down much further with higher orders. That also being said, as per Ch 4 in the book we have a bias-variance trade off to consider. The higher the fit, the more susceptible it will be to variance and thus will not generalise well. I think a linear model is appropriate here given the data.